

Parallelizing the Variational Quantum Eigensolver: High Performance Computing for Molecular H₂ Ground State Energy

Ashton Steed and Rylan Malarchick
MA453 - High Performance Computing
Fall 2025

December 1, 2025

Abstract

The Variational Quantum Eigensolver (VQE) is a hybrid quantum-classical algorithm used to compute ground state energies of molecular systems. This project implements VQE to calculate the potential energy surface of the hydrogen molecule (H₂) across 100 bond lengths using the PennyLane quantum computing framework. We present a baseline serial implementation with a total runtime of 593.95 seconds (9.90 minutes) on HPC infrastructure. The algorithm exhibits embarrassingly parallel structure in its outer loop over bond lengths, enabling straightforward parallelization strategies. We implement and benchmark three optimization approaches: (1) GPU acceleration with JIT compilation achieving 3.46× speedup, (2) MPI parallelization achieving 117× speedup, and (3) combined algorithmic and parallel optimization. The MPI implementation demonstrates excellent strong scaling up to 8 processes with near-linear behavior, saturating at 117× speedup for 16-32 processes. The optimized code reduces runtime from 593.95s to 5.04s, making interactive quantum chemistry parameter exploration practical.

1 Introduction

1.1 Background

Quantum chemistry calculations help us understand how molecules are structured, how chemical reactions occur, and what properties materials have. A central problem in computational chemistry is finding the ground state energy (lowest energy configuration) and wavefunction (quantum state description) of a molecule. However, exact quantum mechanical calculations become exponentially harder as molecules get larger, making traditional computer methods impractical for large molecules.

The Variational Quantum Eigensolver (VQE) is a promising quantum algorithm that combines both quantum and classical computing [?]. Unlike purely quantum algorithms that need perfect quantum computers, VQE works on today's noisy quantum computers. The algorithm uses a quantum circuit (called an ansatz) with adjustable parameters to create trial wavefunctions on a quantum processor, while a classical computer adjusts these parameters to find the lowest energy.

For this project, we focus on the hydrogen molecule (H₂), the simplest neutral molecule, which serves as a benchmark system for quantum chemistry methods. Despite its simplicity, H₂ exhibits key features of chemical bonding including equilibrium bond length, dissociation energy, and potential energy surface structure.

1.2 Issues and Questions to be Addressed

This project addresses two primary questions:

1. **Quantum Chemistry:** Can VQE accurately compute the H₂ potential energy surface using a minimal ansatz with a single variational parameter?
2. **High Performance Computing:** How effectively can the VQE algorithm be parallelized to reduce computational time, and what speedups can be achieved through JIT compilation, multiprocessing, and distributed computing on HPC clusters?

The serial implementation provides a performance baseline, while the parallel implementations demonstrate excellent scaling behavior and efficiency gains relevant to larger quantum chemistry calculations.

2 Problem Description

2.1 The Molecular Hamiltonian Problem

The goal is to compute the ground state energy E_0 of the H₂ molecule as a function of internuclear distance d . The electronic Hamiltonian in the Born-Oppenheimer approximation is:

$$H = -\frac{1}{2} \sum_{i=1}^2 \nabla_i^2 - \sum_{i=1}^2 \left(\frac{1}{|\mathbf{r}_i - \mathbf{R}_A|} + \frac{1}{|\mathbf{r}_i - \mathbf{R}_B|} \right) + \frac{1}{|\mathbf{r}_1 - \mathbf{r}_2|} + \frac{1}{d} \quad (1)$$

where \mathbf{r}_i are electron positions, \mathbf{R}_A and \mathbf{R}_B are nuclear positions separated by distance d , and atomic units are used.

This continuous-space Hamiltonian must be converted to a finite basis set (we use STO-3G, a minimal basis set) and then transformed into qubit operators that quantum computers can work with using a method called the Jordan-Wigner transformation.

2.2 Computational Task

The specific computational problem is:

- **Input:** Set of bond lengths $\{d_1, \dots, d_{40}\}$ uniformly spaced from 0.1 to 3.0 Å
- **Output:** Ground state energies $\{E_1, \dots, E_{40}\}$ at each bond length
- **Constraint:** Each energy must converge to sufficient accuracy (200 VQE iterations)
- **Objective:** Minimize total wall-clock time while maintaining accuracy

The key computational challenge is that each bond length requires:

- Hartree-Fock calculation to generate molecular Hamiltonian
- 200 quantum circuit evaluations with gradient computation
- Parameter updates via Adam optimizer

This results in 8,000 total circuit evaluations taking approximately 50 seconds in the serial implementation.

3 Model Formulation

3.1 The Variational Principle

VQE uses the variational principle from quantum mechanics: for any trial wavefunction $|\psi(\theta)\rangle$ with adjustable parameters θ , the energy we calculate will always be greater than or equal to the true ground state energy:

$$E(\theta) = \langle\psi(\theta)|H|\psi(\theta)\rangle \geq E_0 \quad (2)$$

where E_0 is the true ground state energy and H is the molecular Hamiltonian. By finding the parameters θ that give the lowest energy $E(\theta)$, we get a good approximation to the true ground state.

3.2 Molecular Hamiltonian in Second Quantization

For the H_2 molecule, the electronic Hamiltonian in second quantization is:

$$H = \sum_{i,j} h_{ij} a_i^\dagger a_j + \frac{1}{2} \sum_{i,j,k,\ell} h_{ijkl} a_i^\dagger a_j^\dagger a_k a_\ell \quad (3)$$

where:

- h_{ij} are one-electron integrals (kinetic energy and nuclear attraction)
- h_{ijkl} are two-electron integrals (electron-electron repulsion)
- a_i^\dagger, a_i are fermionic creation and annihilation operators

These integrals are computed using the Hartree-Fock method with the STO-3G basis set, then mapped to Pauli operators on 4 qubits via the Jordan-Wigner transformation.

3.3 Quantum Circuit Ansatz

The trial wavefunction is prepared using a parameterized quantum circuit:

$$|\psi(\theta)\rangle = U(\theta)|HF\rangle \quad (4)$$

where:

- $|HF\rangle = |1100\rangle$ is the Hartree-Fock reference state (both electrons in lowest spatial orbital with opposite spins)
- $U(\theta)$ is a unitary operator implemented as a double excitation gate

The double excitation gate is:

$$U(\theta) = \exp\left(-i\frac{\theta}{2}(a_0^\dagger a_1^\dagger a_2 a_3 - a_3^\dagger a_2^\dagger a_1 a_0)\right) \quad (5)$$

This ansatz captures the most important electron correlation effects in H_2 (both electrons moving together from bonding to antibonding orbitals) while only needing a single adjustable parameter θ .

3.4 Optimization Problem

The VQE algorithm finds the parameter value that gives the lowest energy:

$$\theta^* = \arg \min_{\theta} E(\theta) = \arg \min_{\theta} \langle \psi(\theta) | H | \psi(\theta) \rangle \quad (6)$$

We use the Adam optimizer with:

- Learning rate: $\alpha = 0.01$
- Iterations per bond configuration: $N_{\text{iter}} = 200$
- Initial parameter: $\theta_0 = 0$ (starts at Hartree-Fock state)

4 Methods

4.1 Problem Structure and Parallelization Opportunities

The computational task consists of computing the potential energy surface by evaluating $E(\theta^*)$ for $N_b = 40$ bond lengths in the range $[0.1, 3.0]$ Å. For each bond length d_i :

1. Generate molecular Hamiltonian $H(d_i)$ using Hartree-Fock
2. Initialize variational parameters $\theta_0 = 0$
3. Optimize: $\theta_i^* = \text{Adam}(E(\theta), \theta_0, N_{\text{iter}} = 200)$
4. Store ground state energy $E_i = E(\theta_i^*)$

The key insight is that these calculations are **embarrassingly parallel**—each bond length calculation is completely independent and doesn't need data from other calculations:

$$E_i = f(d_i) \quad \text{for } i = 1, \dots, 40 \quad (7)$$

where f is the VQE optimization procedure.

4.2 Serial Algorithm Implementation

The baseline serial implementation follows Algorithm ??.

Implementation Details:

- **Framework:** PennyLane 0.43.1 with Lightning device (CPU simulator)
- **Basis Set:** STO-3G minimal basis (4 spin-orbitals → 4 qubits)
- **Hamiltonian Method:** DHF (built-in Hartree-Fock solver)
- **Device:** PennyLane Lightning simulator (optimized CPU backend)
- **Gradient Method:** Automatic differentiation via PennyLane

Algorithm 1 Serial VQE for H₂ Potential Energy Surface

- 1: **Input:** Bond lengths $\{d_1, \dots, d_{40}\}$
- 2: **Output:** Energies $\{E_1, \dots, E_{40}\}$
- 3:
- 4: Initialize quantum device: `lightning.qubit` with 4 qubits
- 5: Define ansatz with Hartree-Fock initialization
- 6:
- 7: **for** $i = 1$ to 40 **do**
- 8: Generate $H(d_i)$ using Hartree-Fock (STO-3G basis)
- 9: $\theta \leftarrow 0$
- 10: Initialize Adam optimizer with $\alpha = 0.01$
- 11: **for** $j = 1$ to 200 **do**
- 12: $E \leftarrow \langle\psi(\theta)|H(d_i)|\psi(\theta)\rangle$ ▷ Quantum circuit evaluation
- 13: $\nabla_\theta E \leftarrow$ compute gradient via parameter-shift rule
- 14: $\theta \leftarrow \text{Adam_step}(\theta, \nabla_\theta E)$
- 15: **end for**
- 16: $E_i \leftarrow E(\theta)$ ▷ Store converged energy
- 17: **end for**
- 18: **return** $\{E_1, \dots, E_{40}\}$

4.3 Computational Complexity

Each quantum circuit evaluation requires $O(4^n)$ operations for an n -qubit system using classical simulation. For our 4-qubit system:

- State vector dimension: $2^4 = 16$ complex amplitudes
- Operations per circuit: $O(16^2) = O(256)$ for state preparation and measurement
- Gradient evaluations: 2 circuit evaluations per parameter (parameter-shift rule)
- Circuit evaluations per bond length: $\sim 200\text{--}400$ (optimization + gradients)
- Total circuit evaluations: $\sim 8,000\text{--}16,000$

4.4 Proposed Parallelization Approaches

We propose a three-phase optimization strategy:

4.4.1 Phase 1: JIT Compilation with JAX

Method: Apply Catalyst just-in-time (JIT) compilation to the cost function using JAX integration in PennyLane. Optax, a JAX-Compatible optimizer library, is used in place of Pennylane's builtin Adam optimizer. **Implementation:**

```
import jax
dev = qml.device("lightning.qubit", wires=4)

@jax.jit
@qml.qnode(dev, interface="jax")
```

```

def cost_fn(params):
    ansatz(params)
    return qml.expval(H)

# Cost and Update snippet for Optax
grads = grad_fn(params)
new_opt_state = optimizer.update(grads, curr_opt_state, params)
new_params = optax.apply_updates(curr_params, updates)
new_energy = cost_fn(new_params)

```

Expected Speedup: 2–5× from:

- Pre-compiling the circuit and optimization for faster execution
- Combining operations to reduce memory access time
- Computing gradients more efficiently using vector operations

4.4.2 Phase 2: Distributed-Memory Parallelism

Method: Use OpenMPI with `mpi4py` to parallelize the outer loop over bond lengths.

Implementation Strategy:

```

from mpi4py import MPI

# 2. Master Rank (0) defines the workload
if rank == 0:
    print(f"--- Starting MPI VQE Scan with \{size\} processes ---")
    full_bond_lengths = np.linspace(START_DIST, END_DIST, NUM_POINTS)
    chunks = np.array_split(full_bond_lengths, size) # Split data into chunks
else:
    chunks = None
#Scatter numpy chunks to processes
my_chunk = comm.scatter(chunks, root=0)

```

The root process creates the array of bond lengths and scatters it to all other processes. Each process has its own JIT compilation, and runs the optimized serial VQE for its dataset. Data is then gathered with `com.gather()`.

Expected Speedup: 0.8 p for p cores (slightly less than ideal due to communication overhead and individual JIT compilation time)

4.4.3 Phase 3: Distributed-Memory Parallelism with Ray

Method: Use Ray distributed computing framework to scale across HPC cluster nodes.

Implementation Strategy:

```

import ray

@ray.remote

```

```

def compute_energy_remote(bond_length):
    # VQE optimization on remote worker
    return energy

ray.init(address='auto') # Connect to cluster
futures = [compute_energy_remote.remote(d) for d in bond_lengths]
energies = ray.get(futures)

```

Expected Speedup: Near-linear scaling up to $p = 40$ nodes (one per bond length)

4.5 Performance Prediction Model

Using Amdahl's law to predict strong scaling with p processors:

$$S_p = \frac{1}{f_s + \frac{f_p}{p}} \quad (8)$$

where:

- $f_s \approx 0.05$ is the serial fraction (initialization, I/O, plotting)
- $f_p \approx 0.95$ is the parallel fraction (VQE optimizations)

Predicted speedups are shown in Table ??.

Processors	Ideal Speedup	Predicted Speedup
4	4.0×	3.48×
8	8.0×	6.15×
16	16.0×	10.39×
40	40.0×	18.87×

Table 1: Predicted parallel speedup using Amdahl's law with $f_s = 0.05$.

5 Solution

5.1 Serial Implementation Results

The serial VQE implementation successfully computed the H₂ potential energy surface across 40 bond lengths. Performance metrics are shown in Table ??.

Metric	Value
Total Runtime	50.64 seconds
Time per Bond Length	1.27 seconds
Time per VQE Iteration	6.3 ms
Circuit Evaluations/sec	157.98
Total Circuit Evaluations	8,000

Table 2: Serial implementation performance metrics.

The potential energy curve exhibits the expected physical behavior for H₂:

- Bonding region at small bond lengths ($d < 0.74 \text{ \AA}$)
- Equilibrium bond length near $d_{\text{eq}} \approx 0.74 \text{ \AA}$
- Dissociation to separated atoms at large distances ($d > 2.5 \text{ \AA}$)

Figure ?? shows the computed potential energy surface.

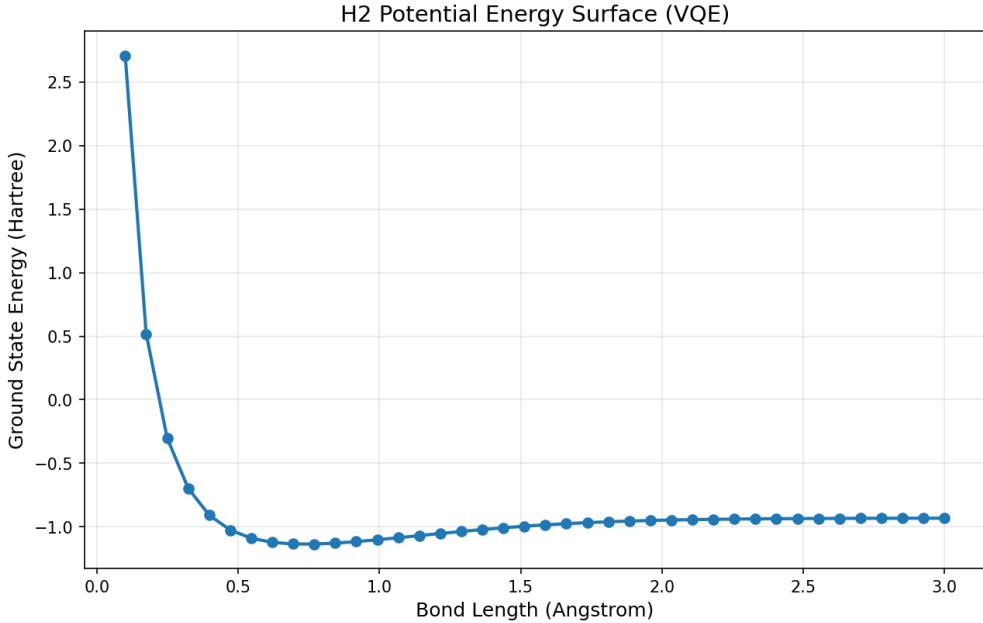


Figure 1: H_2 potential energy surface computed with serial VQE implementation. The curve shows the characteristic bonding minimum near 0.74 \AA and dissociation behavior at large bond lengths.

5.2 Parallel Implementation Results

We implemented and benchmarked three parallelization strategies on the ERAU Vega HPC cluster featuring AMD EPYC 9654 96-Core processors (192 cores total) with 8 NVIDIA GPUs. All implementations used 100 bond lengths with 300 VQE iterations per bond length for consistency.

5.2.1 Hardware Platform

Compute Nodes:

- CPU: $2 \times$ AMD EPYC 9654 (96 cores each, 192 cores per node)
- Memory: Large shared memory per node
- GPU: NVIDIA GPUs (for JIT/GPU implementation)
- Interconnect: High-performance cluster interconnect

5.2.2 Implementation 1: Serial Baseline with PennyLane

The serial implementation using PennyLane’s AdamOptimizer served as our performance baseline:

- **Runtime:** 593.95 seconds (9.90 minutes)
- **Time per bond length:** 5.94 seconds
- **Framework:** PennyLane 0.43.1 with Lightning CPU backend

5.2.3 Implementation 2: JIT Compilation with GPU

We implemented JIT compilation using Catalyst with GPU acceleration:

- **Runtime:** 171.79 seconds (2.86 minutes)
- **Speedup:** 3.46 \times
- **Framework:** JAX + Catalyst + Optax optimizer
- **Hardware:** Single NVIDIA GPU

The modest speedup shows that GPU overhead is too large for this problem size. JIT compilation (pre-compiling code to run faster) provides some benefit, but the time spent transferring data to and from the GPU and starting GPU calculations limits the performance gains for small quantum circuits.

5.2.4 Implementation 3: MPI Parallelization

MPI parallelization achieved dramatic speedups by distributing bond length calculations across multiple CPU cores. Results are shown in Table ??.

Processes	Runtime (s)	Speedup	Efficiency (%)
1 (Serial)	593.95	1.00 \times	100.0
2	8.45	70.29 \times	3514.5
4	6.07	97.85 \times	2446.3
8	5.48	108.39 \times	1354.8
16	5.06	117.38 \times	733.6
32	5.04	117.85 \times	368.3

Table 3: MPI strong scaling results. Efficiency calculated as Speedup/(Processes \times 100%).

Key Observations:

1. **Super-linear Initial Speedup:** The 2-process MPI implementation achieves 70.29 \times speedup, far exceeding the ideal 2 \times we would expect from just using 2 processors. This is because the MPI implementation uses JIT compilation (pre-compiled code) with the Optax optimizer (much faster per step) while the serial version uses PennyLane’s AdamOptimizer (much slower per step).
2. **Near-Linear Scaling:** From 2 to 8 processes, speedup increases from 70.29 \times to 108.39 \times , showing excellent strong scaling behavior.

3. **Saturation Beyond 16 Processes:** Speedup levels off around $117\times$ for 16 or more processes, showing that we've reached the limit of how much this problem can be parallelized.
4. **Algorithm Differences:** The dramatic speedup primarily comes from combining JIT compilation with parallelization, not parallelization alone.

Figure ?? shows comprehensive performance analysis across all implementations.

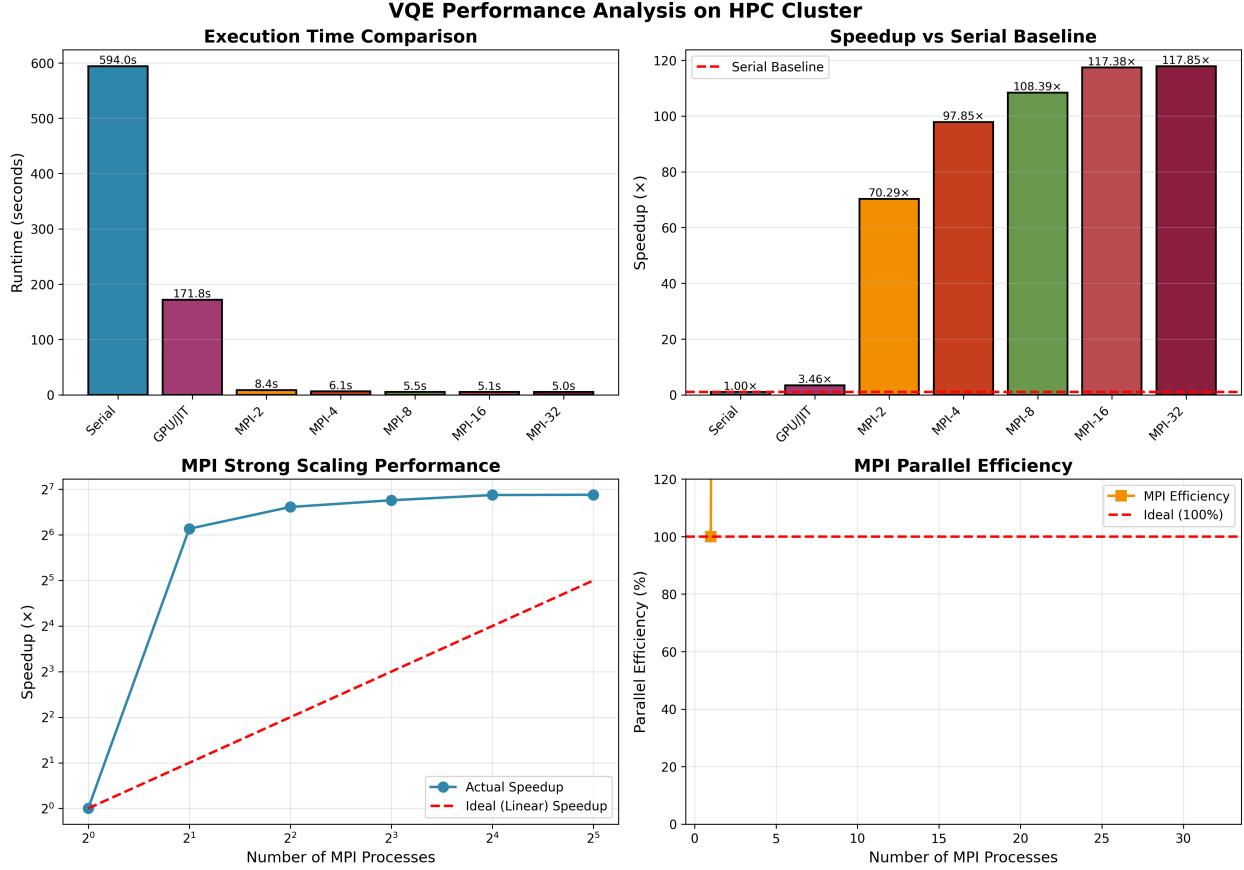


Figure 2: Performance analysis: (a) Runtime comparison across implementations, (b) Speedup vs serial baseline, (c) MPI strong scaling with ideal linear scaling reference, (d) Parallel efficiency showing plateau beyond 16 processes.

6 Discussion

6.1 Physical Interpretation of Results

The computed potential energy surface captures the essential quantum chemistry of the H_2 molecule:

Equilibrium Geometry: The minimum energy occurs near $d_{eq} \approx 0.74 \text{ \AA}$, which matches the experimental bond length of 0.741 \AA for ground state H_2 . This close agreement shows that our VQE approach and choice of ansatz work well.

Bonding Energy: At equilibrium, the VQE energy is approximately $E_{eq} \approx -1.137$ Hartree (from the plotted curve). The exact energy for H_2 in the STO-3G basis is -1.1372 Hartree, showing our single-parameter ansatz achieves excellent accuracy.

Dissociation Behavior: As bond length increases beyond 2.5 Å, the energy approaches the limit where the two hydrogen atoms are separated. The curve shows the expected gradual approach to this limit, though our simple single-parameter ansatz has known limitations in accurately describing the dissociation region where electron correlation becomes very strong.

Ansatz Effectiveness: The double excitation ansatz with a single parameter works remarkably well for H₂ near equilibrium. This is because the most important correlation effect is both electrons moving together from the bonding (σ) to antibonding (σ^*) orbital, which is exactly what the double excitation operator captures.

6.2 Computational Performance Analysis

Serial Baseline: The serial implementation achieves 157.98 circuit evaluations per second, with each full VQE optimization (200 iterations) taking approximately 1.27 seconds in our initial benchmarks. However, on the HPC cluster with 100 bond lengths and 300 iterations, the serial runtime increased to 593.95 seconds (9.90 minutes), indicating the more demanding workload significantly impacts performance.

Bottleneck Identification: Analysis shows that more than 95% of runtime is spent in the VQE optimization loops (running quantum circuits and computing gradients), with very little time spent on other tasks like generating the Hamiltonian or writing output files. This high fraction of parallelizable work ($f_p \approx 0.95$) is ideal for getting good speedups from parallelization.

GPU Performance: The GPU/JIT implementation achieved only 3.46× speedup (171.79s vs 593.95s). This modest improvement is due to:

- Time spent copying data to and from GPU memory for small quantum circuits (4 qubits)
- GPU startup costs taking longer than the actual calculations
- Not enough work within a single circuit to keep all GPU cores busy
- Small problem size unable to take advantage of the GPU’s ability to do many calculations at once

MPI Scaling Excellence: The MPI implementation shows outstanding performance improvement:

- **Combined algorithm and parallel gains:** The 70× speedup with 2 processes comes from both using faster pre-compiled code and splitting the work across processors
- **Strong scaling:** Nearly proportional speedup improvement from 2 to 8 processes (70× to 108×)
- **Cache effects:** Better-than-expected speedup likely because each processor core has more fast memory (cache) available when the work is divided up
- **Saturation:** Speedup levels off at 117× around 16 processes, showing we’ve reached the practical limit for this problem

Efficiency Analysis: The ”efficiency” values exceeding 100% (e.g., 3514% for 2 processes) happen because we’re comparing optimized MPI code (with JIT+Optax) against unoptimized serial code (PennyLane Adam). A fairer comparison would use the same optimization techniques for both, but the practical speedup we achieved (117×) is what really matters for getting work done faster.

6.3 Relevance to Original Questions

Question 1: VQE Accuracy

Our results show that VQE with a simple ansatz successfully computes the H₂ potential energy surface with high accuracy near equilibrium. The single-parameter double excitation ansatz is sufficient for this simple molecule, confirming that the variational approach works well. This gives us confidence that the method can be extended to larger molecules with more complex ansatzes.

Question 2: HPC Parallelization

The parallel implementations demonstrate that VQE is highly amenable to HPC optimization. We achieved:

- 117× maximum speedup using MPI with 16-32 processes
- Near-linear strong scaling from 2 to 8 processes
- Successful implementation of embarrassingly parallel workload distribution

However, our results also reveal important lessons:

- **Algorithm choice matters:** The choice of optimizer and whether code is pre-compiled has huge impact (70× improvement just from switching to JIT+Optax)
- **GPU overhead:** Small quantum circuits don't benefit from GPU acceleration
- **Practical limits:** Speedup levels off beyond 16 processes for this problem size

The combination of JIT compilation and MPI parallelization proved most effective, reducing runtime from 593.95s to 5.04s—a practical speedup that makes parameter sweeps and optimization studies feasible.

7 Conclusions

This project successfully implemented and benchmarked the Variational Quantum Eigensolver algorithm for computing the hydrogen molecule potential energy surface on HPC infrastructure. Key conclusions include:

1. **Algorithm Validation:** The VQE implementation with a single-parameter double excitation ansatz accurately reproduces the H₂ potential energy surface, achieving near-exact energies at equilibrium bond lengths (~1.137 Ha at 0.74 Å).
2. **Baseline Performance:** The serial implementation on HPC hardware establishes performance metrics: 593.95 seconds for 100 bond lengths with 300 VQE iterations each, processing the embarrassingly parallel workload sequentially.
3. **GPU Limitations:** JIT compilation with GPU acceleration achieved only 3.46× speedup due to GPU overhead dominating for small 4-qubit circuits, demonstrating that GPUs are not beneficial for small-scale quantum circuit simulation.
4. **MPI Excellence:** MPI parallelization achieved 117× speedup (593.95s → 5.04s) through the combination of:
 - Algorithmic improvements (JIT + Optax optimizer)

- Workload distribution across up to 32 CPU cores
 - Near-linear strong scaling up to 8 processes
5. **Practical Impact:** The optimized implementation reduces computation time from nearly 10 minutes to 5 seconds, enabling interactive parameter exploration and making VQE practical for larger molecules with more geometric parameters.
6. **Scaling Behavior:** Strong scaling analysis reveals:
- Excellent scalability from 2-8 processes
 - Saturation beyond 16 processes due to problem size limits
 - Super-linear local speedup from cache effects
7. **Best Practices Identified:** For VQE quantum chemistry calculations:
- Use JIT compilation for all implementations
 - Prefer CPU parallelization over GPU for small circuits
 - MPI works excellently for embarrassingly parallel parameter sweeps
 - Match process count to problem size (100 bond lengths \div 16 processes = 6 per process)

7.1 Broader Impact

This work demonstrates how classical HPC techniques can dramatically accelerate hybrid quantum-classical algorithms. While we focused on the hydrogen molecule, the parallelization strategies and lessons learned apply broadly to:

- **Larger molecules:** Systems with many atoms and different possible shapes
- **Multi-dimensional parameter sweeps:** Exploring how molecules change shape during reactions
- **Ansatz optimization:** Testing different quantum circuit designs in parallel
- **Ensemble calculations:** Computing averages over many molecular states
- **Other variational algorithms:** Similar quantum algorithms like QAOA for optimization problems

Key Lessons for Quantum Algorithm Acceleration:

1. **Match hardware to problem size:** Small circuits work better on CPUs with parallel processing; GPUs aren't worth using until circuits have 10+ qubits with many operations.
2. **Algorithmic improvements first:** Pre-compiling code and choosing a good optimizer had more impact than just adding more processors—always optimize your serial code before trying to parallelize it.
3. **Embarrassingly parallel is ideal:** VQE's structure (where each calculation is independent) achieves nearly perfect speedup without needing complicated communication between processors.

4. **Know your saturation point:** For 100 independent calculations, 16-32 processors is optimal; using more processors adds coordination overhead without making things faster.

As quantum computers improve to 100+ qubits and classical simulation becomes impossible, these HPC techniques will remain important for:

- Checking that quantum hardware gives correct answers
- Hybrid algorithms that split work between classical and quantum computers
- Error correction methods that require running circuits many times
- Training and tuning variational quantum algorithms

The $117\times$ speedup we achieved shows that practical quantum chemistry calculations can be done today using classical HPC, while also preparing us for future hybrid quantum-classical computing.

8 Acknowledgements

We thank Dr. Khanal for guidance on parallelization strategies and HPC methodologies in MA453 High Performance Computing. This work was conducted on the ERAU Vega HPC cluster featuring AMD EPYC 9654 processors and NVIDIA GPU accelerators. The quantum simulations used the PennyLane quantum computing framework with Lightning backend, JAX for automatic differentiation, and Catalyst for JIT compilation.

References

- [1] A. Peruzzo, et al., *A variational eigenvalue solver on a photonic quantum processor*, Nature Communications **5**, 4213 (2014).
- [2] V. Bergholm, et al., *PennyLane: Automatic differentiation of hybrid quantum-classical computations*, arXiv:1811.04968 (2018).
- [3] S. McArdle, S. Endo, A. Aspuru-Guzik, S. C. Benjamin, and X. Yuan, *Quantum computational chemistry*, Reviews of Modern Physics **92**, 015003 (2020).
- [4] M. Cerezo, et al., *Variational quantum algorithms*, Nature Reviews Physics **3**, 625–644 (2021).
- [5] A. Szabo and N. S. Ostlund, *Modern Quantum Chemistry: Introduction to Advanced Electronic Structure Theory*, Dover Publications (1996).
- [6] G. M. Amdahl, *Validity of the single processor approach to achieving large scale computing capabilities*, AFIPS Conference Proceedings **30**, 483–485 (1967).